

# PA02: Network I/O Data Movement Cost Analysis

**Student:** Arjun Prasad

**Roll No:** MT25064

**System:** Intel ThinkCentre M70s Gen 3 (Ubuntu 22.04.5 LTS)

## Introduction

This report presents an experimental analysis of data movement costs in TCP socket communication by comparing three implementations:

**A1:** Two-copy baseline using `send()`

**A2:** One-copy optimization using `sendmsg()` with scatter-gather I/O

**A3:** Zero-copy implementation using `MSG_ZEROCOPY`

The experiments measured throughput, CPU cycles, cache misses, and context switches across varying message sizes (64B - 64KB) and thread counts (1-8).

## Part A: Implementation Details

### A1. Two-Copy Implementation

#### Q. Where Do the Two Copies Occur?

In the standard `send()` implementation, data undergoes at least two copy operations:

1. **First Copy (User → Kernel):**

- Application calls `send()`
- Kernel copies data from user-space buffer to kernel socket buffer.
- This copy is performed by the kernel via `copy_from_user()`

2. **Second Copy (Kernel → Network Stack):**

- Data is copied from socket buffer to DMA buffer.
- Network card copies data from system memory to its own buffers

#### Q. Is it actually only two copies?

There may be additional copies due to the following reasons:-

**Third copy:** If TCP segmentation is needed, data may be copied again.

**Cache line copies:** CPU cache operations involve implicit copies between cache levels.

**DMA transfers:** While DMA is often considered “zero-copy,” it still moves data.

### Code Implementation

```
// Eight separate send() calls for the 8-field message structure
send(sock, msg->field1, field_size, 0);
send(sock, msg->field2, field_size, 0);
// ... (8 total sends)
```

Each `send()` incurs a system call overhead and separate copy operation.

### A2. One-Copy Implementation

#### Q. How is One Copy Eliminated?

The `sendmsg()` implementation with scatter-gather I/O (`iovec`) eliminates the user-space consolidation copy:

#### Traditional approach - 2 copies:

1. User allocates large buffer.
2. Copies 8 fields into contiguous buffer.
3. Calls send() on consolidated buffer.

#### **1 copy - approach:**

1. User keeps 8 fields in separate malloc'd buffers.
2. Creates iovec array pointing to each buffer.
3. Calls sendmsg() once with iovec array.
4. Kernel performs scatter-gather DMA directly from the 8 buffers.

#### **Code implementation:**

##### **Without sendmsg:**

```
char *consolidated = malloc(total_size);
memcpy(consolidated, field1, size);      // Copy 1
memcpy(consolidated + size, field2, size); // Copy 2
// ...
send(sock, consolidated, total_size, 0); // Copy to kernel
Total: 8 memcpy + 1 kernel copy = 9 operations
```


##### **With sendmsg:**


```
struct iovec iov[8];
iov[0].iov_base = field1; iov[0].iov_len = size;
iov[1].iov_base = field2; iov[1].iov_len = size;
// ...
sendmsg(sock, &msghdr, 0); // Kernel gathers directly
Total: 1 kernel gather operation = 1 operation
```

The kernel's scatter-gather engine can read from multiple discontiguous memory regions directly, eliminating the user-space memcpy operations.

### A3. Zero-Copy Implementation

#### *Kernel Behavior Diagram*





### Code Implementation:

```

// Page-aligned allocation
char *buf;
posix_memalign((void**)&buf, 4096, msg_size);

// Enable zero-copy on socket
int val = 1;
setsockopt(sock, SOL_SOCKET, SO_ZEROCOPY, &val, sizeof(val));

// Send with MSG_ZEROCOPY flag
sendmsg(sock, &msghdr, MSG_ZEROCOPY);

```

## Part B: Experimental Results

### System Configuration

- **CPU:** Intel Core (ThinkCentre M70s Gen 3)
- **OS:** Ubuntu 22.04.5 LTS
- **Kernel:** Linux (perf\_event\_paranoid = -1)
- **Network:** Loopback (127.0.0.1)
- **Measurements:** 48 total experiments (3 implementations × 4 sizes × 4 thread counts)

### Summary Results Table

Implementation	Msg Size	Threads	Throughput (Gbps)	Cycles	Cache Misses	Ctx Switches
A1 (2-copy)	64	1	0.126	8.1B	5.2M	872
A1 (2-copy)	64	8	0.875	259B	3.5M	2,364
A1 (2-copy)	65536	4	233.2	127B	198M	1,933
A2 (1-copy)	64	1	1.209	14B	112M	123
A2 (1-copy)	4096	8	235.4	252B	19M	4,412
A2 (1-copy)	65536	4	286.4	131B	100M	1,996
A3 (0-copy)	64	1	0.261	23.5B	16M	722
A3 (0-copy)	65536	4	231.4	124B	180M	1,587
A3 (0-copy)	65536	8	78.7	24.6B	31M	144,488

## Part C: Automation Script

The bash script (`MT25064_Part_C_run_experiments.sh`) automates the following:


1. **Compilation:** Compiles all three server implementations and client.
2. **Experiment execution:** Runs 48 experiments systematically.
3. **Perf collection:** Captures CPU cycles, cache misses, context switches using `perf stat`.
4. **Data storage:** Saves results in CSV format with encoded parameters.

### Execution:

```
sudo ./MT25064_Part_C_run_experiments.sh
# Outputs: MT25064_Part_C_results.csv
```

## Part D: Visualization and Analysis


Plot 1: Throughput vs Message Size



### Observations:

- A2 dominates across all message sizes**
  - 10× better than A1 at 64 bytes (1.2 vs 0.13 Gbps)
  - Reaches 286 Gbps at 65KB with 4 threads
- Zero-copy underperforms at small messages**
  - A3 marginally better than A1.
  - Notification overhead negates benefits
- Throughput scales with message size**
  - Logarithmic relationship visible
  - System call overhead amortized over larger messages
- Thread saturation at 8 threads**
  - A2 and A3 show throughput degradation at 65KB
  - Cache contention and context switching overhead


*Plot 2: Throughput vs Thread Count (4096 bytes)*



**Observations:**

1. **Near-linear scaling (1→4 threads)**
  - A2: 37 → 150 Gbps (4× improvement)
  - Good parallelism efficiency
2. **A2 super-linear jump at 8 threads**
  - 150 → 235 Gbps (1.57× from 4→8 threads)
  - Suggests CPU pipeline optimizations kicking in
3. **A1 and A3 converge at high thread counts**
  - Both reach ~78-79 Gbps at 8 threads
  - Copy overhead becomes less significant vs. scheduling overhead
4. **Performance ranking consistent:**
  - A2 > A3 ≥ A1 at all thread counts

*Plot 3: Cache Misses vs Message Size*




Cache Misses vs Message Size

#### Observations:

1. **A1 spike at 512 bytes**
  - 111 million cache misses (worst case!)
  - Likely hitting cache line boundary inefficiencies
  - Message size causes poor alignment with 64-byte cache lines
2. **A3 shows best cache behavior at medium sizes**
  - Drops from 64M → 18M misses (64 → 4096 bytes)
  - Zero-copy eliminates buffer copy cache pollution
  - Direct DMA from user buffers = better locality
3. **Large message cache misses increase**
  - 65KB messages cannot fit in L3 cache
  - All implementations show ~50-100M misses
  - Capacity misses dominate
4. **A2 shows most stable cache behavior**
  - Gradual, predictable increase
  - Scatter-gather minimizes intermediate buffers

**Winner:** A3 at medium sizes, A2 for consistency

Plot 4: CPU Cycles per Byte



#### Observations:

- Dramatic efficiency improvement with size**
  - 64 bytes: 50-300 cycles/byte
  - 65536 bytes: 0.2-1 cycles/byte
  - **System call overhead amortization**
- A2 is most efficient**
  - Orange lines consistently lowest
  - 65KB @ 8 threads: 0.04 cycles/byte
  - Scatter-gather DMA is extremely CPU-efficient
- Zero-copy NOT most efficient**
  - A3 uses more cycles/byte than A2 at small sizes
  - Kernel notification mechanism adds overhead
  - Only competitive at very large messages
- Thread scaling helps efficiency**
  - More threads → lower cycles/byte
  - Parallel processing amortizes fixed costs
  - Exception: Contention at 8 threads for some configs

**Efficiency Ranking:** A2 > A1 > A3 (at small sizes)

## Part E: Analysis and Reasoning

### Q1: Why does zero-copy not always give the best throughput?

**Answer:**

Zero-copy (MSG\_ZEROCOPY) does not always win because of kernel notification overhead that outweighs the benefits at small message sizes.

Specific mechanisms causing overhead:

#### 1. Page Pinning Cost:

- Kernel must pin user pages in physical memory using `get_user_pages()`
- Prevents page swapping, adds TLB pressure
- For small messages, pinning overhead > copy cost

#### 2. Reference Counting:

- Kernel maintains reference counts on user pages
- Must track when DMA completes before unpinning
- Atomic operations on reference counts cause cache coherency traffic

#### 3. Completion Notification:

- Application must poll error queue with `recvmsg(MSG_ERRQUEUE)`
- Adds system call overhead
- For small messages, notification cost > copy cost

#### 4. Context Switches:

- Our data shows A3 with 8 threads @ 65KB: 144,488 context switches
- A2 same config: 28,372 context switches
- Zero-copy notification mechanism causes severe scheduling overhead

**Conclusion:** Zero-copy requires large messages (>4KB) and moderate parallelism to be beneficial.

### Q2: Which cache level shows the most reduction in misses and why?

**Answer:**

The Last Level Cache (LLC/L3) shows the most reduction in misses, as measured by perf's `cache-misses` event, which specifically tracks LLC misses.

**Evidence from observation:**

*Comparing A1 vs A2 at 2 threads:*

Message Size	A1 LLC Misses	A2 LLC Misses	Reduction
64 bytes	20.1M	17.4M	13%
512 bytes	111.2M	36.9M	67%
4096 bytes	68.8M	40.1M	42%
65536 bytes	101.3M	45.7M	55%

**Why LLC shows the most reduction:**

#### 1. Working Set Size:

- L1/L2 caches (32KB/256KB) are too small for network buffers
- LLC (8-16MB) can hold multiple message buffers

- Eliminating intermediate copies keeps more data in LLC
- Temporal Locality:**
    - Two-copy: Data touches L1 → L2 → L3 → RAM → L3 → L2 → L1 (multiple evictions)
    - One-copy: Data path is shorter therefore better LLC residency
  - Scatter-Gather Efficiency:**
    - A2's iovec keeps 8 separate buffers in cache
    - No consolidation copy means no cache pollution
    - LLC can hold all 8 buffers simultaneously
  - Cache Line Alignment:**
    - At 512 bytes, A1 shows 111M misses.
    - Likely 512 bytes = 8 cache lines = poor alignment
    - A2's scatter-gather avoids this pathological case

**Note:** - L1/L2 misses are not directly measured by our perf events. However, they're implicit in the LLC miss count. (Most L1/L2 misses are satisfied by LLC.)

**Conclusion:** LLC benefits most because it's the last level before expensive DRAM access. Reducing LLC misses has the highest performance impact.

### Q3: How does thread count interact with cache contention?

#### Answer:

Thread count has a non-linear relationship with cache contention, showing good scaling up to 4 threads, then severe degradation at 8 threads due to cache thrashing and false sharing.

#### Evidence from observation:

Threads	A2 @ 65KB Throughput	A2 @ 65KB Cache Misses	Context Switches
1	75.6 Gbps	14.0M	1,238
2	138.0 Gbps	45.7M	765
4	<b>286.4 Gbps</b>	99.5M	1,996
8	107.3 Gbps	<b>460.0M</b>	<b>28,372</b>

#### Key Observations:

- Throughput collapse at 8 threads:**
  - Drops from 286 → 107 Gbps (62% reduction!)
  - Cache misses spike 4.6× (100M → 460M)
  - Context switches spike 14× (2K → 28K)
- Cache Coherency Traffic:**
  - Multiple threads access socket buffers
  - MESI protocol (Modified/Exclusive/Shared/Invalid) causes cache line bouncing
  - Each thread invalidates other threads' cache lines
- False Sharing:**
  - TCP socket structures likely share cache lines
  - Lock contention on socket lock causes cache line ping-pong

- Even if data is independent, metadata sharing causes contention
4. **Memory Bandwidth Saturation:**
- ThinkCentre M70s likely has ~40-50 GB/s memory bandwidth
  - At 286 Gbps ≈ 35 GB/s, we're near saturation
  - 8 threads push past bandwidth limit, causing stalls

### **Thread-Specific Patterns:**

1-2 threads: Good cache sharing, minimal contention  
 4 threads: Optimal point - parallelism without saturation  
 8 threads: Cache thrashing + scheduler overhead

### **Architectural Factors:**

- **Cache hierarchy:** L1/L2 are per-core, L3 is shared
- **Core count:** If system has 4-6 physical cores, 8 threads causes oversubscription
- **Hyperthreading:** Logical cores share L1/L2, causing intra-core contention

### **Mitigation Strategies:**

1. **Thread affinity:** Pin threads to specific cores
2. **Separate socket buffers:** Reduce shared state
3. **Batching:** Process multiple messages per thread to improve cache reuse

**Conclusion:** 4 threads is the sweet spot for this workload on this system. Beyond 4, cache contention and scheduler overhead dominate performance.

## **Q4: At what message size does one-copy outperform two-copy on your system?**

### **Answer:**

One-copy (A2) outperforms two-copy (A1) at all message sizes tested in our experiments.

### **Evidence from observation:**

Message Size	Threads	A1 Throughput	A2 Throughput	A2 Speedup
<b>64 bytes</b>	1	0.126 Gbps	1.209 Gbps	<b>9.6×</b>
<b>64 bytes</b>	8	0.875 Gbps	7.769 Gbps	<b>8.9×</b>
<b>512 bytes</b>	1	1.082 Gbps	8.151 Gbps	<b>7.5×</b>
<b>4096 bytes</b>	4	47.51 Gbps	150.6 Gbps	<b>3.2×</b>
<b>65536 bytes</b>	4	233.2 Gbps	286.4 Gbps	<b>1.2×</b>

### **Why A2 Always Wins:**

1. **No user-space consolidation:**
  - A1 would need to memcpy 8 fields together (not in our implementation, but typical)
  - A2 lets kernel do scatter-gather DMA directly
2. **Single system call:**
  - A1 makes 8 separate send() calls (one per field)
  - A2 makes 1 sendmsg() call
  - System call overhead: 8× vs 1×
3. **Better cache behavior:**

- A2 keeps data in separate buffers (better locality)
- A1's multiple sends thrash cache

**Conclusion:** On this system, sendmsg() with iovec should always be preferred over multiple send() calls for structured messages.

## Q5: At what message size does zero-copy outperform two-copy on your system?

### Answer:

Zero-copy (A3) begins to consistently outperform two-copy (A1) at 4096 bytes and above, with optimal performance at moderate thread counts (1-4 threads).

### Evidence:

Message Size	Threads	A1 Throughput	A3 Throughput	Winner	Speedup
64 bytes	1	0.126 Gbps	0.261 Gbps	A3	2.1×
64 bytes	8	0.875 Gbps	1.411 Gbps	A3	1.6×
512 bytes	1	1.082 Gbps	2.073 Gbps	A3	1.9×
512 bytes	8	12.50 Gbps	11.22 Gbps	<b>A1</b>	-
<b>4096 bytes</b>	1	11.96 Gbps	14.35 Gbps	A3	1.2×
<b>4096 bytes</b>	4	47.51 Gbps	54.20 Gbps	A3	1.1×
<b>4096 bytes</b>	8	78.72 Gbps	78.73 Gbps	Tie	1.0×
<b>65536 bytes</b>	1	62.97 Gbps	63.04 Gbps	A3	1.0×
<b>65536 bytes</b>	4	233.2 Gbps	231.4 Gbps	A1	-
<b>65536 bytes</b>	8	170.4 Gbps	78.70 Gbps	<b>A1</b>	-

### Key Findings:

- Transition point: 4096 bytes**
  - Below 4KB: A3 shows marginal gains (1.2-2×)
  - At 4KB: A3 begins to show consistent advantage
  - Above 4KB: Performance depends heavily on thread count
- Thread count critically important:**
  - 1-4 threads: A3 competitive or better
  - 8 threads: A3 collapses (78 vs 170 Gbps at 65KB)
  - Notification overhead scales poorly with threads
- Unexpected behavior at 65KB:**
  - A3 with 8 threads: 144,488 context switches (vs 23,661 for A1)
  - Notification mechanism causes severe scheduler thrashing
  - Zero-copy actually becomes negative-copy in overhead!

### Why 4096 Bytes is the Threshold:

- Page size alignment:**
  - 4096 bytes = 1 memory page
  - DMA engines work optimally at page boundaries
  - Smaller messages waste DMA setup overhead
- Copy cost vs notification cost:**
  - memcpy ~5-10 GB/s on modern CPUs

- 4KB copy:  $\sim 0.4\text{-}0.8 \mu\text{s}$
- MSG\_ZEROCOPY notification:  $\sim 1\text{-}2 \mu\text{s}$
- Crossover at 4KB where copy cost > notification cost

### 3. Cache size considerations:

- 4KB fits in L1 cache (32KB)
- Larger messages spill to L2/L3
- Zero-copy avoids cache pollution at these sizes

## Conclusion:

Zero-copy is beneficial for 4KB-32KB messages with moderate parallelism. Beyond 4 threads or above 32KB, the notification overhead destroys performance. The sweet spot is 4KB-8KB with 2-4 threads, achieving 10-20% improvement over traditional send(). For this workload, `sendmsg()` with iovec (A2) is universally superior to both two-copy and zero-copy.

## Q6: Identify one unexpected result and explain it using OS or hardware concepts.

### Answer:

Throughput degradation at 8 threads with 65KB messages for A2 and A3:

- A2: Drops from 286.4 Gbps (4T) to 107.3 Gbps (8T) - 62% reduction
- A3: Drops from 231.4 Gbps (4T) to 78.7 Gbps (8T) - 66% reduction
- A1: Drops from 233.2 Gbps (4T) to 170.4 Gbps (8T) - 27% reduction

This is unexpected because:

1. More threads should increase throughput (Amdahl's Law)
2. The system has 8+ logical cores available
3. A1 (two-copy) shows much smaller degradation than optimized versions

### Explanation Using OS and Hardware Concepts:

#### 1. Memory Bandwidth Saturation

**Theoretical Analysis:** - ThinkCentre M70s Gen 3: DDR4-3200 ( $\sim 25\text{-}40 \text{ GB/s}$  bandwidth) - Peak throughput at 4T: 286 Gbps = **35.75 GB/s** - At 8T: Attempting to push beyond memory bandwidth ceiling

**Evidence:** - A2 @ 4T: 286 Gbps (sustainable) - A2 @ 8T: 107 Gbps (bandwidth-limited) - System hits **memory wall**

**Why A1 degrades less:** - A1 already slower, doesn't hit bandwidth limit - Operating below saturation point - More headroom for additional threads

#### 2. Cache Coherency Protocol Overhead (MESI/MESIF)

##### Mechanism:

Thread 1 writes → Cache line in "Modified" state

Thread 2 reads → Cache line invalidated, fetched from Thread 1's L1

Thread 3 writes → Entire cache line invalidated across all cores

**With 8 threads:** - **Cache line bouncing** between cores - Each write invalidates 7 other cores' cache lines - Coherency traffic saturates interconnect

##### Evidence from Context Switches:

Implementation	4 Threads	8 Threads	Increase
A1	1,933	23,661	12.2×
A2	1,996	28,372	14.2×
A3	1,587	<b>144,488</b>	<b>91×</b>

A3's 91× increase in context switches indicates **severe scheduler thrashing**.

### 3. False Sharing in Socket Structures

#### TCP Socket Buffer Structure (Simplified):

```
struct sock {
    spinlock_t lock;           // ← Shared across threads
    struct sk_buff_head write_queue; // ← Highly contended
    atomic_t wmem_alloc;        // ← Atomic updates
    // ... (likely 64-128 bytes total)
};
```

**False Sharing Scenario:** - Multiple threads call sendmsg() on same socket - Socket lock at byte offset 0 - Queued packet count at byte offset 8 - Both in same **64-byte cache line** - Thread 1 updates lock → invalidates Thread 2's cache line - Ping-pong effect causes stalls

**Why worse for A2/A3:** - Faster individual operations mean more lock acquisitions/second - A1 is slower, so less contention per unit time - A2/A3 beat themselves by being "too fast"

### 4. CPU Scheduler Overhead

**Linux CFS Scheduler Behavior:** - With 8 threads on ~6-8 cores: Oversubscription - Time slice: ~1-10ms per thread - Context switch cost: ~2-5µs

**At high throughput:** - 286 Gbps = 35.75 GB/s - 65KB messages = ~550,000 messages/second - With 8 threads: ~69,000 messages/thread/second - **1 message every 14µs**

**Critical observation:** - Message processing time: 14µs - Context switch overhead: 2-5µs - **Overhead = 14-35% of work time!**

With 28,000 context switches over 10 seconds: - 2,800 switches/second -  $2,800 \times 3\mu s = 8.4\text{ms lost to switching}$  - At 8.4ms/second = **0.84% overhead** (seems small) But cache effects multiply this: - Each switch flushes L1/L2 caches - Reload time: ~100-500 cycles - With 2.8K switches: ~280M-1.4B wasted cycles - Matches our perf data showing 165-252B total cycles

### 5. Non-Uniform Memory Access (NUMA) Effects

#### If system has NUMA:

CPU 0-3: Memory Controller 0

CPU 4-7: Memory Controller 1

**Remote memory access penalty:** - Local: ~60ns latency - Remote: ~120ns latency (2× penalty)

**With 8 threads:** - Some threads allocated to remote NUMA node - Socket buffers might be on different NUMA node - Cross-node traffic adds latency

**Evidence:** - A2 @ 4T might run on single NUMA node - A2 @ 8T forces cross-NUMA communication - Explains 62% throughput drop

### 6. Kernel Network Stack Lock Contention

#### Linux TCP/IP stack has several global locks:

```
// Simplified kernel code
```

```
socket_send() {
```

```

    spin_lock(&sk->sk_lock);           // Per-socket Lock
    // ... process data ...
    spin_lock(&tcp_write_queue);     // Queue Lock
    // ... enqueue ...
    spin_unlock(&tcp_write_queue);
    spin_unlock(&sk->sk_lock);
}

```

**Lock hold time:** - A1: Longer per-call (more work) = less frequent locking - A2: Shorter per-call (optimized) = MORE lock acquisitions

**Paradox:** - Optimizing the code path makes it **more sensitive to lock contention** - 8 threads hammering same socket = lock convoy

**Solution:** - Use **separate sockets** per thread (not implemented in our design) - Or use SO\_REUSEPORT to distribute across sockets

## AI Usage Declaration

Components Where AI Was Used

### 1. Part A - Implementation

- **Prompts:**
  - Explain how Linux MSG\_ZEROCOPY works internally with page pinning”
  - “What are the differences between sendmsg with iovc vs multiple send calls”
- **Usage:** Clarified kernel implementation details, verified understanding
- **Own work:** All code implementations, struct definitions, explanations

### 2. Part C - Bash Scripting

- **Prompts:**
  - “Why is perf stat not writing output to file”
  - “How to parse CSV output from perf stat -x,”
- **Usage:** Debugging process management, perf output parsing
- **Own work:** Experiment design, loop structure, CSV formatting logic

### 3. Part D - Plotting

- **Prompts:**
  - “How to create log-log plot in matplotlib with multiple lines”
  - “Best way to layout legend outside plot area”
- **Usage:** Plot aesthetics, formatting, legend positioning
- **Own work:** Data hardcoding, plot selection, interpretation

### 4. Part E - Analysis

- **Used:** Claude for brainstorming explanation structure
- **Prompts:**
  - “What are common causes of cache coherency overhead in multithreaded programs”
  - “Explain NUMA effects on network performance”
- **Usage:** Confirming technical concepts, structure for explanations
- **Own work:** All analysis, conclusions, data interpretation, unexpected result identification

## 5. Report Writing

- **Prompts:**
  - “Format this data as a markdown table”
  - “How to create ASCII diagram for kernel data flow”
- **Usage:** Document formatting, diagram syntax
- **AI helped with:** Syntax, formatting, debugging, concept verification

### Tools Used

- **ChatGPT-4** (OpenAI): Concept clarification, debugging
- **Claude-3.5** (Anthropic): Report structuring, technical review

## GitHub Repository

Repository URL: [https://github.com/arjun25064/MT25064-GRS/GRS\\_PA02](https://github.com/arjun25064/MT25064-GRS/GRS_PA02)

## References

1. Linux Kernel Documentation: Documentation/networking/msg\_zerocopy.rst
2. sendmsg(2) man page
3. Intel® 64 and IA-32 Architectures Optimization Reference Manual
4. “The Linux Programming Interface” by Michael Kerrisk (Chapter 61: Sockets)
5. Linux perf documentation: <https://perf.wiki.kernel.org/>
6. [https://linux.how2shout.com/what-is-perf\\_event-in-linux-linux-performance-monitoring/](https://linux.how2shout.com/what-is-perf_event-in-linux-linux-performance-monitoring/)
7. <https://www.brendangregg.com/perf.html>